Algorithms and Datastructures Shortest Path, Dijkstra Algorithm

Albert-Ludwigs-Universität Freiburg

Prof. Dr. Rolf Backofen

Bioinformatics Group / Department of Computer Science Algorithms and Datastructures, March 2018

Structure



Graphs

Dijkstra Algorithm

Structure



Graphs

Dijkstra Algorithm

```
For a graph G = (V, E):
```

- A path of *G* is a sequence of edges $u_1, u_2, ..., u_i \in V$ with
 - Undirected graph: $\{u_1, u_2\}, \{u_2, u_3\}, \dots, \{u_{i-1}, u_i\} \in E$
 - Directed graph: $(u_1, u_2), (u_2, u_3), \dots, (u_{i-1}, u_i) \in E$
- The length of a path is
 - Without weights: number of edges taken
 - With weights: sum of weigths of edges taken

For a graph G = (V, E):

- The shortest path between two vertices u, v is the path P = (u, ..., v) with the shortest length d(u, v) or lowest costs
- The diameter of a graph is the longest shortest path

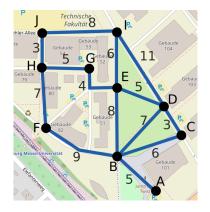
Structure



Graphs

Dijkstra Algorithm

- Wanted: Shortest path from M to all other points
- Place pearls on crossings and clamp strings between them



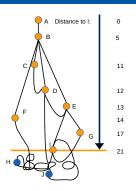
Shortest Path without Computer





Abbildung: Based on OpenStreetMaps; CC BY-SA 2.0

Take the net and pull it slowly upwards until fully lifted



- Each node (pearl) now has a specific height
- The distance to M is exactly the shortest path



Abbildung: Shortest path from s to t

- Let r be the shortest path from s to t
- For each node u on path r the path from u to t is the shortest path

Proof:

- If there was a shorter path from s to u then we could choose this path to get faster to t
- Then *r* would not be the shortest path



Abbildung: Shortest path from s to t

- This is also correct for all sub paths on r
- If the shortest path from s to t passes u_1 and u_2 then the sub path (u_1, u_2) is the shortest path from u_1 to u_2

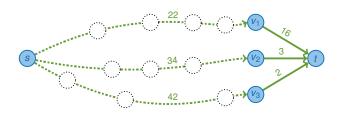


Abbildung: Shortest paths from s to t

■ If we know the shortest path form s to the preceding nodes of $t(v_1, v_2, v_3)$ we can determine the shortest path to t

Idea:

- Attach the cost of the shortest path to each node
- Let the information travel over the edges (message passing)
- In which order should we process the nodes?

Inventor:

- Edsger Dijkstra (1930 2002)
- Computer scientist from Netherlands
- Won Turing-Award as one of few Europeans for his studies of structured programming
- Invented the Dijkstra-Algorithm in 1959



Abbildung: Portrait © Hamilton Richards - manuscripts of Edsger W. Dijkstra, University Texas at Austin

Example:

- Lift pearl *M* a little bit
- Connections to pearls R, L and G are hanging in the air
- Lift further until pearl R starts to lift at 5 m
- The shortest path to R is now known
- Lift further: The wires from R,O and Q are now in the air
- One of the pearls G, L, Q or O is the next one Which one?

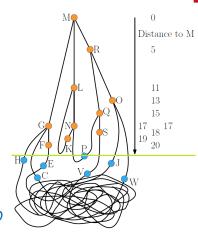


Abbildung: Map © Mehlhorn / Sanders

Example:

- At 11 m pearl *L* gets lifted
- The wires to *N* and *K* are now in the air
- One of the pearls G, K, N, Q or O is the next one Which one?
- At 13m pearl O gets lifted ...
- How to translate this into an computer algorithm?

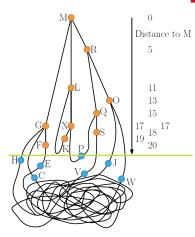


Abbildung: Map © Mehlhorn / Sanders



High level description: Three types of nodes

Settled: For node u we know dist(s, u) (Pearl example: This pearl is hanging in the air)



Active: For node u we know a tentative distance $td(u) \ge dist(s, u)$ (Can be optimal but doesn't have to) (Pearl example: This pearl is laying on the table but one connected wire is already in the air)



Unreached: We have not reached the node yet (Pearl example: This preal is hanging in the air)



High level description:

- Each iteration take the active node u with the smallest td(u) (The pearl getting lifted next)
- We update the state of the node *u* to settled (The pearl gets lifted)
- We check for each neighbor v of node u if we can reach v faster than currently possible (Check all outgoing wires from this pearl: Activate all connected pearls, update tentative distance if smaller)
- Iterate until no active nodes exist anymore



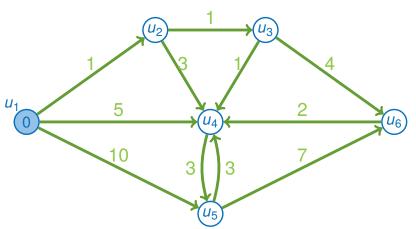
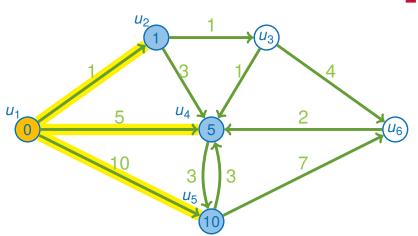
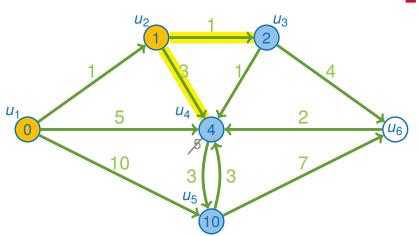


Abbildung: Start at u₁

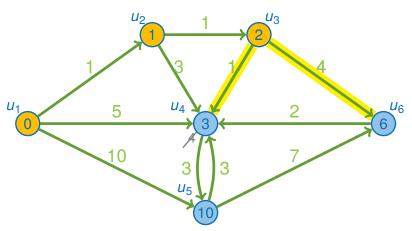




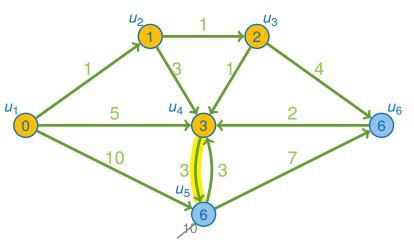




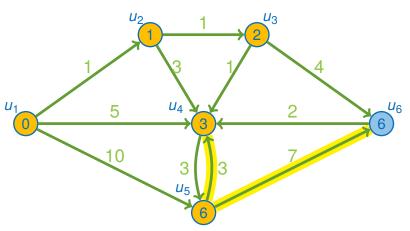




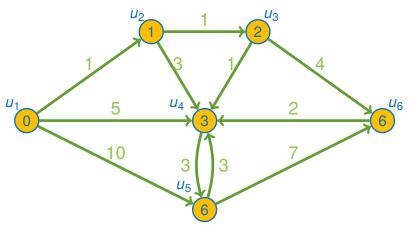














- Assumption 1: All edges have a positive length
- **Assumption 2:** Each node has a unique distance dist(s, u) (This was not the case on the previous slides)

This results in an easy and intuitive proof.

It is prossible to show this without assumption

It is prossible to show this without assumption 2. See references if interested

■ With assumption 2 there exists a sorting $u_1, u_2, ...$ with that:

$$\operatorname{dist}(s, u_1) < \operatorname{dist}(s, u_2) < \operatorname{dist}(s, u_3) < \dots$$

Proof:

With **assumption 2** there exists a sorting $u_1, u_2, ...$ with that:

$$\operatorname{dist}(s, u_1) < \operatorname{dist}(s, u_2) < \operatorname{dist}(s, u_3) < \dots$$

- We want to show that the *Dijkstra* algorithm finds the shortest path for each node u_i so that $td(u_i) = dist(s, u_i)$ holds
- Additionally we show that each node gets solved in order of the distance: Node u_i gets solved in iteration i

$$u_1, u_2, u_3, \dots$$

To show: Node u_i gets solved in round i

- Node u_i contains the correct distance $(td(u_i) = dist(s, u_i))$ and is active
- Node u_i has the smallest value for $td(u_i)$ and gets selected by the algorithm

Induction start:

- Only the start node $s = u_1$ is active and td(s) = 0
 - Node u_1 gets solved and $td(u_1) = dist(s, u_1) = 0$)
- 2 Only the start node u_1 is active

Induction step: i = i + 1

- **To show:** Node u_i contains the correct distance $(td(u_i) = dist(s, u_i))$ and is active
 - \blacksquare On the shortest path from s to u_{i+1} is a preceding node that:

$$\operatorname{dist}(s,u_{i+1})=\operatorname{dist}(s,v)+\operatorname{c}(v,u_{i+1})$$

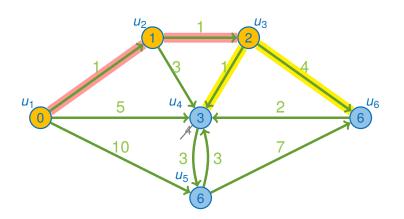
(c are the costs of the edge)



- With that results $dist(s, v) < dist(s, u_{i+1})$ because c > 0
- Because u_{i+1} is currently solved node v is one of the preceding nodes u_1, \ldots, u_i , hence $v = u_i$ with $0 \le j \le i$

Proof - Example of Iteration 6





- Preceding node of u_6 is $v = u_3$
- In round 3 $td(u_6) = 2 + 4 = 6$ was already solved



- **To show:** Node u_i contains the correct distance $(td(u_i) = dist(s, u_i))$ and is active
 - With **induction assumption**: v already contains the correct distance which was evaluated in round j (edge from v to u_{i+1}) and is stored in $td(u_{i+1})$
 - u_{i+1} is active because the preceding node was solved



- **To show:** Node u_{i+1} has the smallest value for $td(u_{i+1})$ and gets selected by the algorithm
 - All nodes with smaller dist are already solved
 - All other nodes u_k with k > i + 1 have a greater $dist(s, u_k)$ an with that the $td(u_k)$ is greater or equal
 - $\Rightarrow u_{i+1}$ is the node with the smallest td and gets selected by the algorithm

Implementation:

Implementation

- We have to manage a set of active nodes
- We start with only the start node in our set
- At the start of each iteration we need the node u with the smallest td(u)

How to implement this?

Implementation:

- Using a priority queue with td(u) as keys
- The following problem occurs:
 - The tentative distance of an active node might change multiple times before it is settled
 - We have to change the key in our priority queue without removing the entry

Limitations:

- Often only insert, getMin and deleteMin are implemented
- ⇒ We only have access to the first element and not any desired one

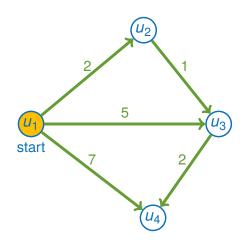
Alternative:

Implementation

- If a node reoccurs with a smaller dist we insert the element one more time into the priority queue (We do nothing if the distance is greater or equal)
- We do not remove the old entry
- The node always gets solved with the smallest distance because of the priority
- If a node reoccurs with a higher dist we remove it and do simply nothing

Implementation - Example

Priority queue:



Graph with *n* nodes and *m* edges: $(m \ge n)$

- Each node gets solved exactly one time
- When solving a node it's outgoing edges are taken into account
- Each edge triggers at maximum one insert operation
- The number of operations on the priority queue is at maximum O(m)
- This results in a runtime of $O(m \cdot \log m)$ (log m because of at max. m elements in the priority queue)

Runtime of $O(m \cdot \log m)$:

- Because of $m < n^2$ we have a maximum runtime of $O(m \cdot \log n)$, because $\log n^2 = 2 \log n$
- With a complex priotity queue the runtime can be reduced to $O(m+n\log n)$
 - For example with a Fibonacci heap
 - This results in a better runtime for complex graphs $m \sim n^2$
 - Complex heaps create a management overhead
 - \Rightarrow In practice $m \in O(n)$ with a **binary heap** being faster (See lecture 6)

Termination criteria:

Terminate as soon as the target node t is settled ... never before because tentative distance might change:

$$td(t) \ge dist(s, t)$$

Before the node t is solved all nodes u with $dist(s, u) \leq dist(s, t)$ are settled

Termination criteria:

- Not only the single source single target shortest path problem is solved by the Dijkstra algorithm but also the single source all targets problem
- This sounds wasteful but there is not a (much) better method for general graphs Intuitive: We only know that there is no shorter path if all in the range of dist(s,t) around s is evaluated

- With the current implementation of the Dijkstra algorithm we only get the length of the path How to get the path too?
- If we save the preceding node of the current shortest path on relaxation of each node we can reconstruct the path



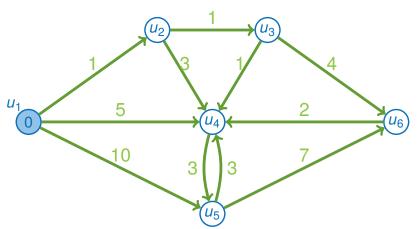
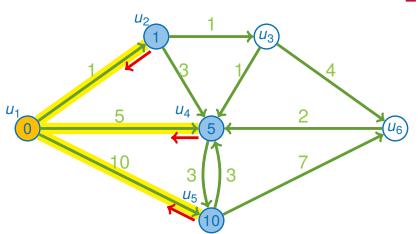
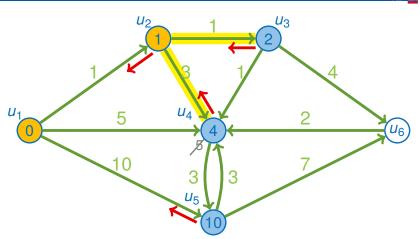


Abbildung: Start at u_1

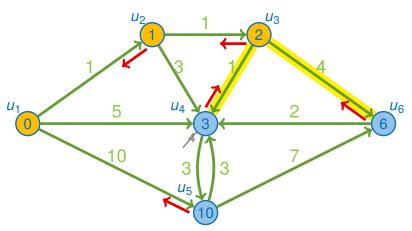




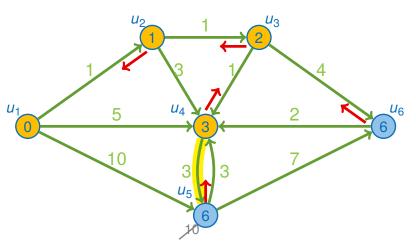




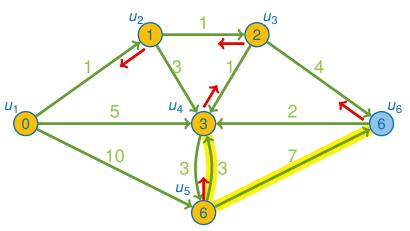




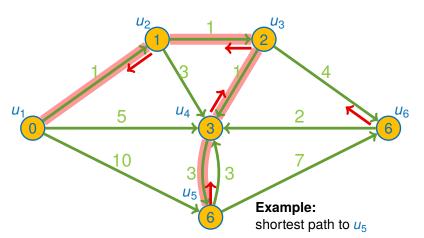












Enhancement:

- In our proof we used the assumption that all costs are not negative (even > 0)
- With negative costs there might be negative cycles:

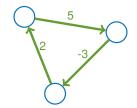


Abbildung: Here no problem ...

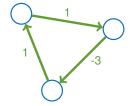
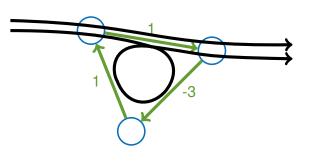


Abbildung: ... but here

Negative cycles:

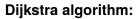


- No cycle: cost of 1
- 1 cycle: cost of 0
- 2 cycles: cost of -1
- 3 cycles:
 - cost of -2
 -

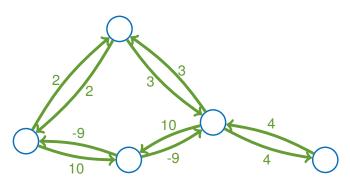
Enhancement:

- We need a different algorithm to deal with negative edges
 - For example the **Bellman-Ford** algorithm
 - If the graph is acyclic we can simply use a topological sorting (with DFS) and ralaxing the nodes in order of this sorting
- Another (not only) in artificial intelligence used variant of the Dijkstra algorithm is the A* algorithm Additional information given:

h(u) = estimated value for dist(u,t)

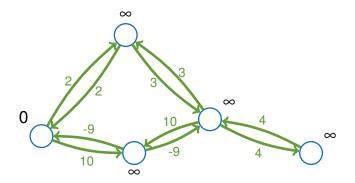


Message passing only from solved nodes



Bellman-Ford algorithm:

Message passing from all nodes until the path lengths are stable



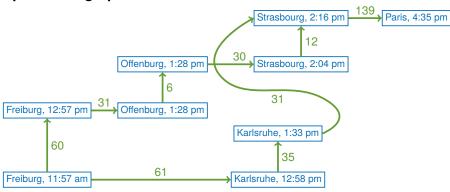
Application example:

- Route planner for car trips (exercise sheet)
- Route planner for bus / train connections What could the graph look like?

Application



Space-time graph:



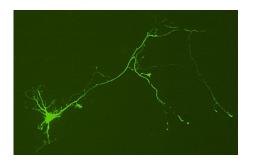


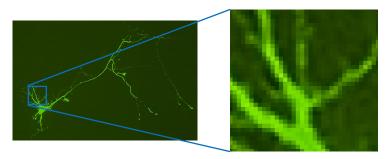


Abbildung: Neurons under fluorescence microscope

- **Task:** Measure length of axons (connections of neurons)
- Demo with ImageJ plugin NeuronJ http://www.imagescience.org/meijering/software/ neuronj/

Application: Trace axons





- Image as graph: Each pixel is a node
- Implicit edges: Each pixel has an edge to it's 8 neighbours (no need to save the edges)
- Costs for nodes (not edges): bright pixels are cheap, dark pixels are costly

■ General

- [CRL01] Thomas H. Cormen, Ronald L. Rivest, and Charles E. Leiserson.

 Introduction to Algorithms.
 - MIT Press, Cambridge, Mass, 2001.
- [MS08] Kurt Mehlhorn and Peter Sanders.

 Algorithms and data structures, 2008.

 https://people.mpi-inf.mpg.de/~mehlhorn/
 ftp/Mehlhorn-Sanders-Toolbox.pdf.

■ Dijkstra's algorithm

[Wik] Dijkstra's algorithm

https:

//en.wikipedia.org/wiki/Dijkstra's_algorithm

- Shortest path problem
 - [Wik] Shortest path problem

https://en.wikipedia.org/wiki/Shortest_path_problem